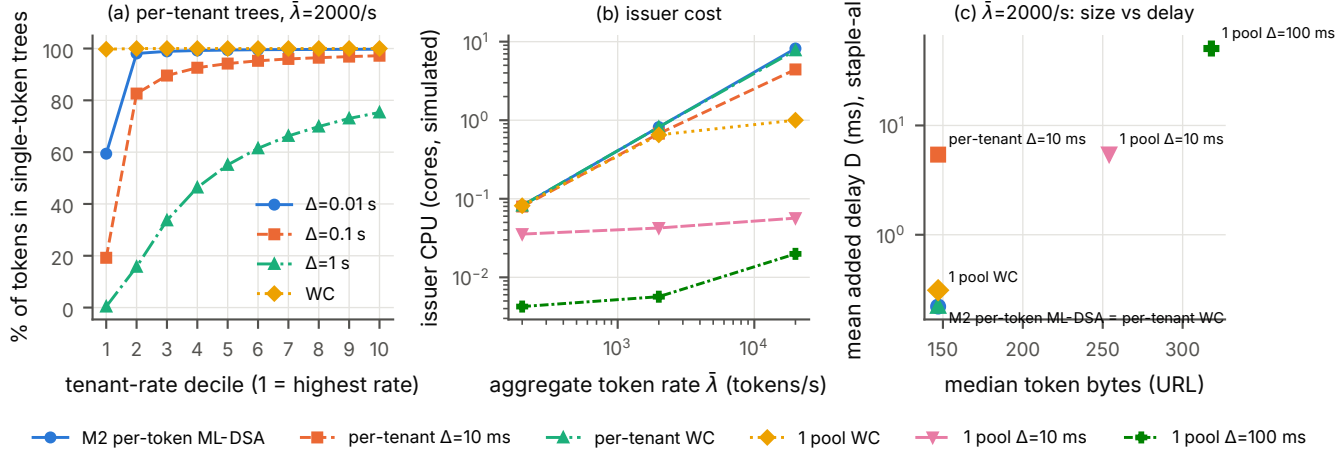
(a) per-tenant trees, $\bar{\lambda}$=2000/s
% of tokens in single-token trees
tenant-rate decile (1 = highest rate)
Δ=0.01 s
Δ=0.1 s
Δ=1 s
WC
(b) issuer cost
issuer CPU (cores, simulated)
aggregate token rate $\bar{\lambda}$ (tokens/s)
(c) $\bar{\lambda}$=2000/s: size vs delay
mean added delay D (ms), staple-all
median token bytes (URL)
per-tenant Δ=10 ms
1 pool Δ=10 ms
1 pool Δ=100 ms
1 pool WC
M2 per-token ML-DSA = per-tenant WC
M2 per-token ML-DSA
per-tenant Δ=10 ms
per-tenant WC
1 pool WC
1 pool Δ=10 ms
1 pool Δ=100 ms